

Lab in Justice Data Science

## Assignment 1

Due 4:10p Wednesday 1/31/2024

Download the file ‘stops\_austin\_tx.zip’ from the Assignment folder on Courseworks, and extract the .csv file inside. These data regard police traffic stops in Austin, Texas from 2005-2016. The dataset is published by the [Stanford Open Policing Project](#). The .csv file is a very small subset of the available data from this organization.

The questions for you to answer are **indicated by bold text**. *Keep your answers short and succinct!* Code here is highlighted in **blue**. You do not need to include any R code when you submit.

1. Read the data file into RStudio using your method of choice (try the ‘Import Dataset’ dropdown in the Environment pane, if you like). Give the dataframe a short but informative name, such as `stops`. **How many rows are in the dataframe? How many columns? What does each row represent?** (Check the [About page online](#) if you’re unsure)
  2. All data is represented in R as a “type”, such as numeric, character, or logical. **What are the data types contained in the following columns: `subject_age`, `search_basis`, `vehicle_model`, `frisk_performed`?** There are multiple ways to get this information, but **find a way to do it with a SINGLE line of code, or NO code, and tell me how you did it** (Try searching Google for “r check type of dataframe column” if you get stuck).
  3. **What percent of this dataset contain occurrences where contraband goods were found?** Use the column `contraband_found` for this, and recall that you access a column in a dataframe using the `$` symbol, like so: `stops$contraband_found`. A powerful feature of logical data is that R treats `TRUE` as `1` and `FALSE` as `0`. This means you can use the `sum()` function to count how many `TRUE`’s are in a column. However, if missing values (`NA`: not available) are present, `sum()` will output something you may not expect (try it). Read the docs by running `?sum` to find an optional argument to `sum()` that solves this problem. Another hint: the function `nrow()` may be helpful here.
  4. Let’s visualize the ages of people who are pulled over. Use the `hist()` function to plot a histogram of the `subject_age` column, and **include the plot here**. **Is there anything difficult to interpret about the axes? How old were the youngest and oldest people**

**stopped?** The `max()` and `min()` functions will help you, and if you get a strange output, see the hint in #3 above.

5. The appearance of a histogram depends on the size of the bins (the width of the bars). We can change this parameter with a different value for the `breaks` argument in the `hist()` function. Read the docs by running `?hist` and **report the default value for `breaks` by looking in the Usage section.**
6. Replace the value for `breaks` by adding `breaks = seq(1,120,10)` to your call to `hist()`. Be sure to separate all function arguments (the inputs inside the parentheses) with a comma followed by a space, and make sure to include the `breaks =` part. **What happened to your histogram when you included this argument? What does the `seq()` function do?** You can either look at the docs, or just type `seq(1,120,10)` into the console and look at the output.
7. Play with the third argument to `seq()` to change the look of your histogram. Try `breaks = seq(1,120,5)`, and `breaks = seq(1,120,2)`. **Which do you think is most informative? Are 16-year-olds pulled over more than 25-year-olds?** *Keyboard shortcut alert!* Press the up arrow key when at the prompt in the console to quickly enter your last line of executed code. You can keep pressing up to move back in your command history. This is much better than typing a function over and over while you're testing – try it!
8. **Make a pie chart to quickly inspect the racial breakdown of traffic stops using the `subject_race` column, and include the plot here.** Remember the `table()` function returns counts of a categorical variable, and the `pie()` function makes a plot, taking in as an argument the results of `table()`. See the R code from Lab 1 on Courseworks for a refresher.
9. The raw counts from the `table()` function aren't particularly informative. We can convert the counts to percentages using math operations, like we could with any vector (or matrix). Assign the results of `table()` to a new object `counts`. Next, divide `counts` by the

sum of `counts` and multiply by 100 to convert the table into percentages. Finally, [check out the federal census data for Austin, TX in 2010](#). **Does there seem to be any racial disparity in who was stopped relative to their representation in the population? Why or why not (use numbers to justify)?** (We would certainly do a more careful analysis before coming to any conclusions)

10. The `table()` function can also make so-called contingency tables, where counts are grouped based on more than one variable. Pass in two columns of the dataset as arguments to `table()`, first `subject_race`, then `subject_sex`. Assign the output to a new variable and convert to percentages, just as you did above. **Does there seem to be any disparity in the sex recorded for people stopped? If so, is it the same for all races? Use numbers to (briefly) justify your answer.**

Finally, I want the lab assignments to be valuable for you – they should be a challenge, but not impossible, nor should they take an egregious amount of time. The questions below are optional, but the more data I have on how you’re doing, the more useful I can make the assignments. Thanks in advance!

**On a scale of 1-5, with 1 being “too easy”, 5 being “too hard”, and 3 being “a good level”, how difficult was this assignment?**

**Approximately how much time did you spend working on this assignment? (i.e. actively solving the problem, not exploring the data independently)**

**Any thoughts on what you found useful, or what you found mundane, or confusing? (anything, big or small)**